

DMRG Tutorial

James Vance

github.com/jnvance

jvance@ictp.it · jvance@up.edu.ph



DMRG.x is a **distributed-memory** implementation of the **Density Matrix Renormalization Group** (DMRG) algorithm based on the **PETSc** and **SLEPc** libraries

Documentation:

<https://DMRGx.rtfd.io>

Source Code:

<https://github.com/jnvance/DMRG.x>

Tutorial Resources:

<https://github.com/jnvance/DMRGTutorial>

Contents of this tutorial

- DMRG algorithm
- Intro to parallel programming with MPI
- PETSc and SLEPc libraries
- Using the DMRG.x application

DMRG Algorithm

Traditional DMRG and its computational aspects

Main Reference:

Feiguin, Adrian E. "The Density Matrix Renormalization Group and its time-dependent variants." *AIP Conference Proceedings*. Vol. 1419. No. 1. AIP, 2011. [\[link\]](#)

Density Matrix Renormalization Group

- ◎ DMRG is one of the most powerful numerical algorithms used to study the ground state properties of many-body systems
- ◎ Developed in 1992 by Steven White to overcome problems in applying NRG to quantum lattice systems
- ◎ Currently being used in many problems in physics and quantum chemistry

Advantages of using DMRG

- ◎ Can treat large systems with high precision with less computational resources than exact diagonalization
- ◎ Unlike quantum Monte Carlo methods, it does not suffer from the negative sign problem in frustrated and fermionic systems

The Infinite-size DMRG Algorithm

Given two blocks and all of its operators

length
↓
 N

basis size
↓
 m

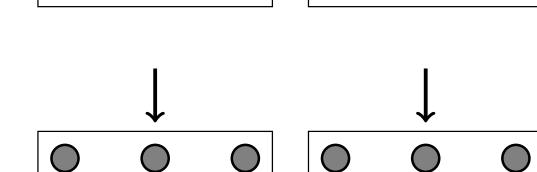
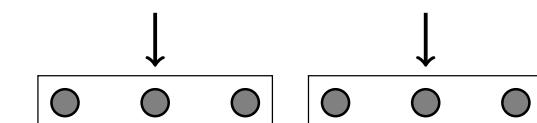
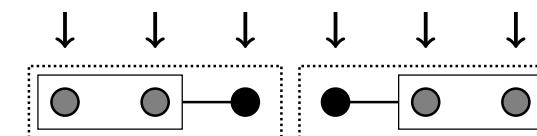
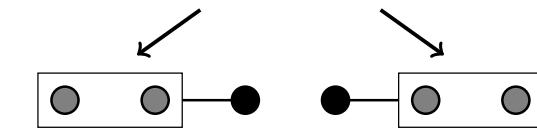
Single DMRG Step

Grow the blocks by **adding one site** (dimension d)

Form the superblock and **diagonalize** the Hamiltonian

From the ground state, determine the **reduced density matrices** (RDMs) of each block

Truncate the basis using the top D_R eigenstates of the RDM



$N + 1$ md

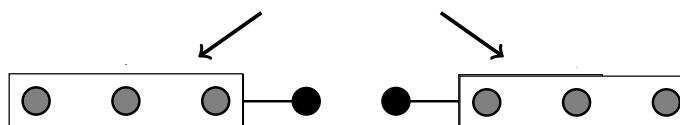
$2N + 2$ $m^2 d^2$

$N + 1$ md

Reduced basis
↓

$N + 1$ m

⋮



$N + 2$ md

The Infinite-size DMRG Algorithm

From a computational perspective

length
↓
 N

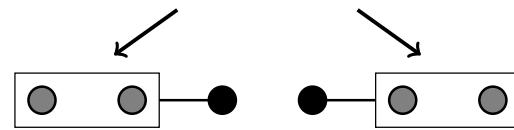
basis size
↓
 m

Setting up a **matrix representation**



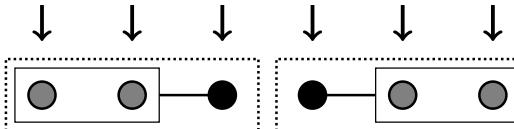
Single DMRG Step

Implementation of a **Kronecker product** routine



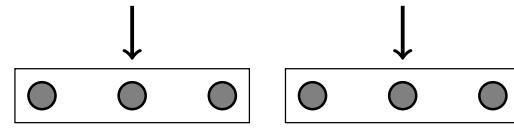
$N + 1$ md

Diagonalization of a large sparse matrix



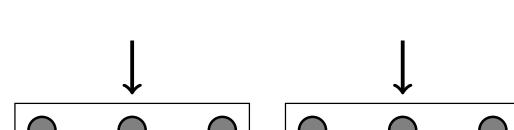
$2N + 2$ m^2d^2

Construction and **eigendecomposition** of the reduced density matrices



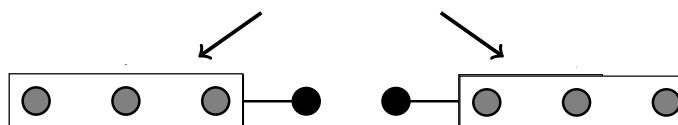
$N + 1$ md

Rotation of the operator matrices to a truncated basis



Reduced basis
↓
 m

⋮

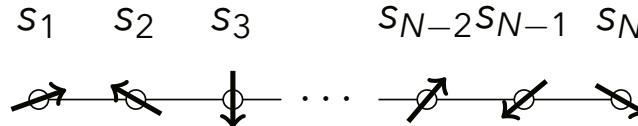


$N + 2$ md

Matrix Representation

Illustrated with the Spin-1/2 Heisenberg Hamiltonian

1D spin chain:



Heisenberg Hamiltonian:

$$\hat{H} = \sum_{i=1}^{N-1} \hat{\mathbf{S}}_i \cdot \hat{\mathbf{S}}_{i+1}$$

$$\hat{H} = \sum_{i=1}^{N-1} \hat{S}_i^z \hat{S}_{i+1}^z + \frac{1}{2} [\hat{S}_i^+ \hat{S}_{i+1}^- + \hat{S}_i^- \hat{S}_{i+1}^+]$$

where

$$S^z = \begin{pmatrix} 1/2 & 0 \\ 0 & -1/2 \end{pmatrix}, S^+ = \begin{pmatrix} 0 & 1 \\ 0 & 0 \end{pmatrix}, S^- = \begin{pmatrix} 0 & 0 \\ 1 & 0 \end{pmatrix},$$

Matrix Representation

Illustrated with the Spin-1/2 Heisenberg Hamiltonian



Two spins:



$$\hat{H} = \hat{S}_1^z \hat{S}_2^z + \frac{1}{2} [\hat{S}_1^+ \hat{S}_2^- + \hat{S}_1^- \hat{S}_2^+]$$

$$H_2 = S^z \otimes S^z + \frac{1}{2} [S^+ \otimes S^- + S^- \otimes S^+]$$



Adding a single site:



$$H_3 = H_2 \otimes \mathbb{1}_2 + \tilde{S}_2^z \otimes S^z + \frac{1}{2} [\tilde{S}_2^+ \otimes S^- + \tilde{S}_2^- \otimes S^+]$$

$$\tilde{S}_2^z = \mathbb{1}_2 \otimes S^z, \quad \tilde{S}_2^\pm = \mathbb{1}_2 \otimes S^\pm$$

Matrix Representation

Illustrated with the Spin-1/2 Heisenberg Hamiltonian



General recursion for adding an i^{th} site to a block
of $i - 1$ sites



$$H_i = H_{i-1} \otimes \mathbb{1}_2 + \tilde{S}_{i-1}^z \otimes S^z + \frac{1}{2} [\tilde{S}_{i-1}^+ \otimes S^- + \tilde{S}_{i-1}^- \otimes S^+]$$

where

$$H_2 = S^z \otimes S^z + \frac{1}{2} [S^+ \otimes S^- + S^- \otimes S^+]$$

$$\tilde{S}_{i-1}^z = \mathbb{1}_{2^{i-2}} \otimes S^z, \quad \tilde{S}_{i-1}^\pm = \mathbb{1}_{2^{i-2}} \otimes S^\pm$$

Matrix Representation of the superblock Hamiltonian



Superblock:



$$\hat{H} = \hat{H}_{L,i+1} + \hat{H}_{R,i+2} + S_{i+1}^z S_{i+2}^z + \frac{1}{2} (S_{i+1}^+ S_{i+2}^- + S_{i+1}^- S_{i+2}^+) \quad \text{---} \quad \text{---} \quad \text{---} \quad \text{---}$$



Explicitly:

$$\begin{aligned} H &= \underbrace{H_{L,i+1} \otimes \mathbb{1}_{D_R \times 2}}_{\text{---}} + \underbrace{\mathbb{1}_{D_L \times 2} \otimes H_{R,i+2}}_{\text{---}} \\ &+ \tilde{S}_{L,i+1}^z \otimes \tilde{S}_{R,i+2}^z \\ &+ \frac{1}{2} \tilde{S}_{L,i+1}^+ \otimes \tilde{S}_{R,i+2}^- \\ &+ \frac{1}{2} \tilde{S}_{L,i+1}^- \otimes \tilde{S}_{R,i+2}^+ \end{aligned} \quad \text{---}$$

Diagonalization of the superblock Hamiltonian

- ◎ Solve the ground state of the superblock Hamiltonian:

$$H |\Psi\rangle = E |\Psi\rangle$$

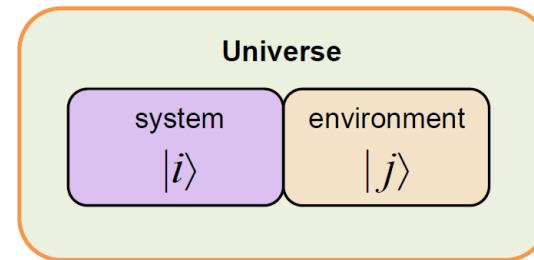
- ◎ Iterative diagonalization using the **Lanczos algorithm**:

$$|\phi_{n+1}\rangle = \boxed{\hat{H}|\phi_n\rangle} - a_n|\phi_n\rangle - b_n^2|\phi_{n-1}\rangle,$$

$$a_n = \frac{\langle\phi_n|\hat{H}|\phi_n\rangle}{\langle\phi_n|\phi_n\rangle}, \quad b_n^2 = \frac{\langle\phi_n|\phi_n\rangle}{\langle\phi_{n-1}|\phi_{n-1}\rangle},$$

- ◎ Only the **matrix-vector product** has to be known
- ◎ The large sparse matrix (size $D_L D_R d^2 \times D_L D_R d^2$)
should not be constructed (**matrix-free approach**)

Construction and Eigendecomposition of the Reduced Density Matrices



- ◎ Bipartite system:
- ◎ Reduced density matrix of the “system” block

$$\hat{\rho}_A = \text{Tr}_B |\Psi\rangle\langle\Psi|$$

$$\rho_{Aii'} = \langle i | \hat{\rho}_A | i' \rangle = \sum_j \langle ij | \Psi \rangle \langle \Psi | i' j \rangle = \sum_j \Psi_{ij} \Psi_{i'j}^*$$

- ◎ Properties of $\hat{\rho}_A$
 1. Hermitian
 2. Non-negative eigenvalues
 3. Unit-trace
 4. Eigenvectors form an orthonormal basis

New eigenvector basis:

$$\hat{\rho}_A = \sum_{\alpha} \omega_{\alpha} |\alpha\rangle\langle\alpha|;$$

with $\omega_{\alpha} \geq 0$ and $\sum_{\alpha} \omega_{\alpha} = 1$

Density Matrix Truncation and Rotation to the New Basis

- ◎ Eigendecomposition of the reduced density matrix

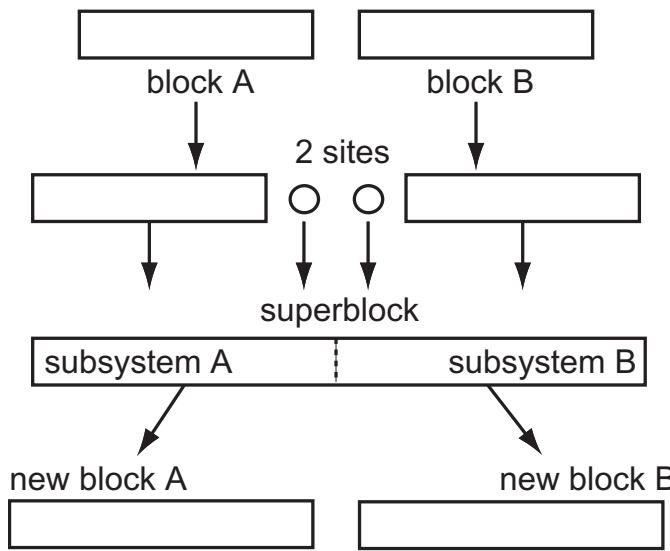
$$\hat{\rho}_A = \sum_{\alpha} \omega_{\alpha} |\alpha\rangle\langle\alpha|$$

- ◎ Keep at most m states with the largest eigenvalues ω_{α} forming a rectangular rotation operator U ($md \times m$)
- ◎ Rotate each block operator O_A to the truncated basis

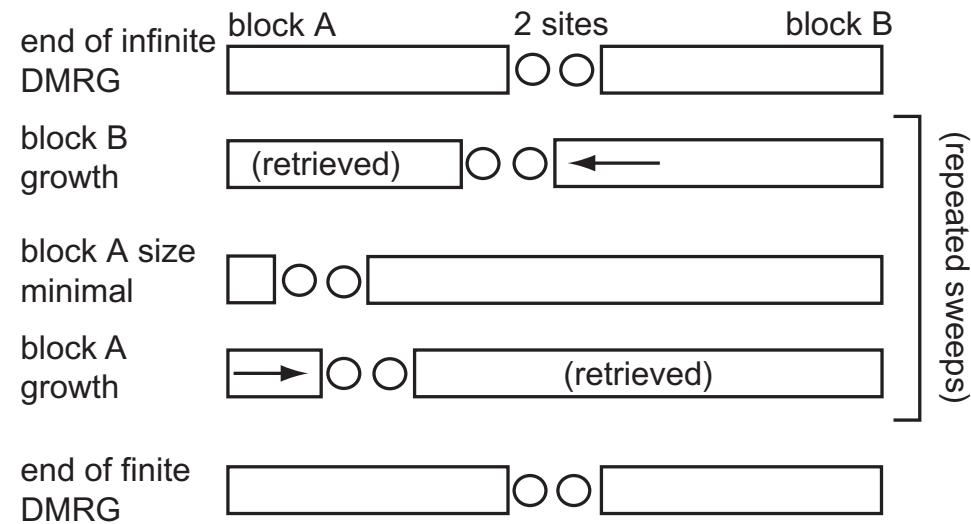
$$\tilde{O}_A = U^{\dagger} O_A U$$

Diagram illustrating the rotation of block operator O_A :
Input: $md \times md$ (represented by a red box)
Output: $m \times m$ (represented by a red box)
The transformation is shown as $U^{\dagger} O_A U$, where U^{\dagger} is the conjugate transpose of the rotation operator U .

DMRG Algorithm



Infinite-size algorithm (warmup)



Finite-size algorithm (sweeps)

Sample codes and links

<https://github.com/jnvance/DMRGTutorial>

To download: **Clone or download**

- git clone <url>
- Download ZIP

SIMPLE-DMRG

Prototype Implementation in Python

Main Reference:

James R. Garrison, & Ryan V. Mishmash. (2017, November 29). simple-dmrg/simple-dmrg: Simple DMRG 1.0 (Version v1.0.0). Zenodo. <http://doi.org/10.5281/zenodo.1068359>

Prototype Implementation: SIMPLE-DMRG

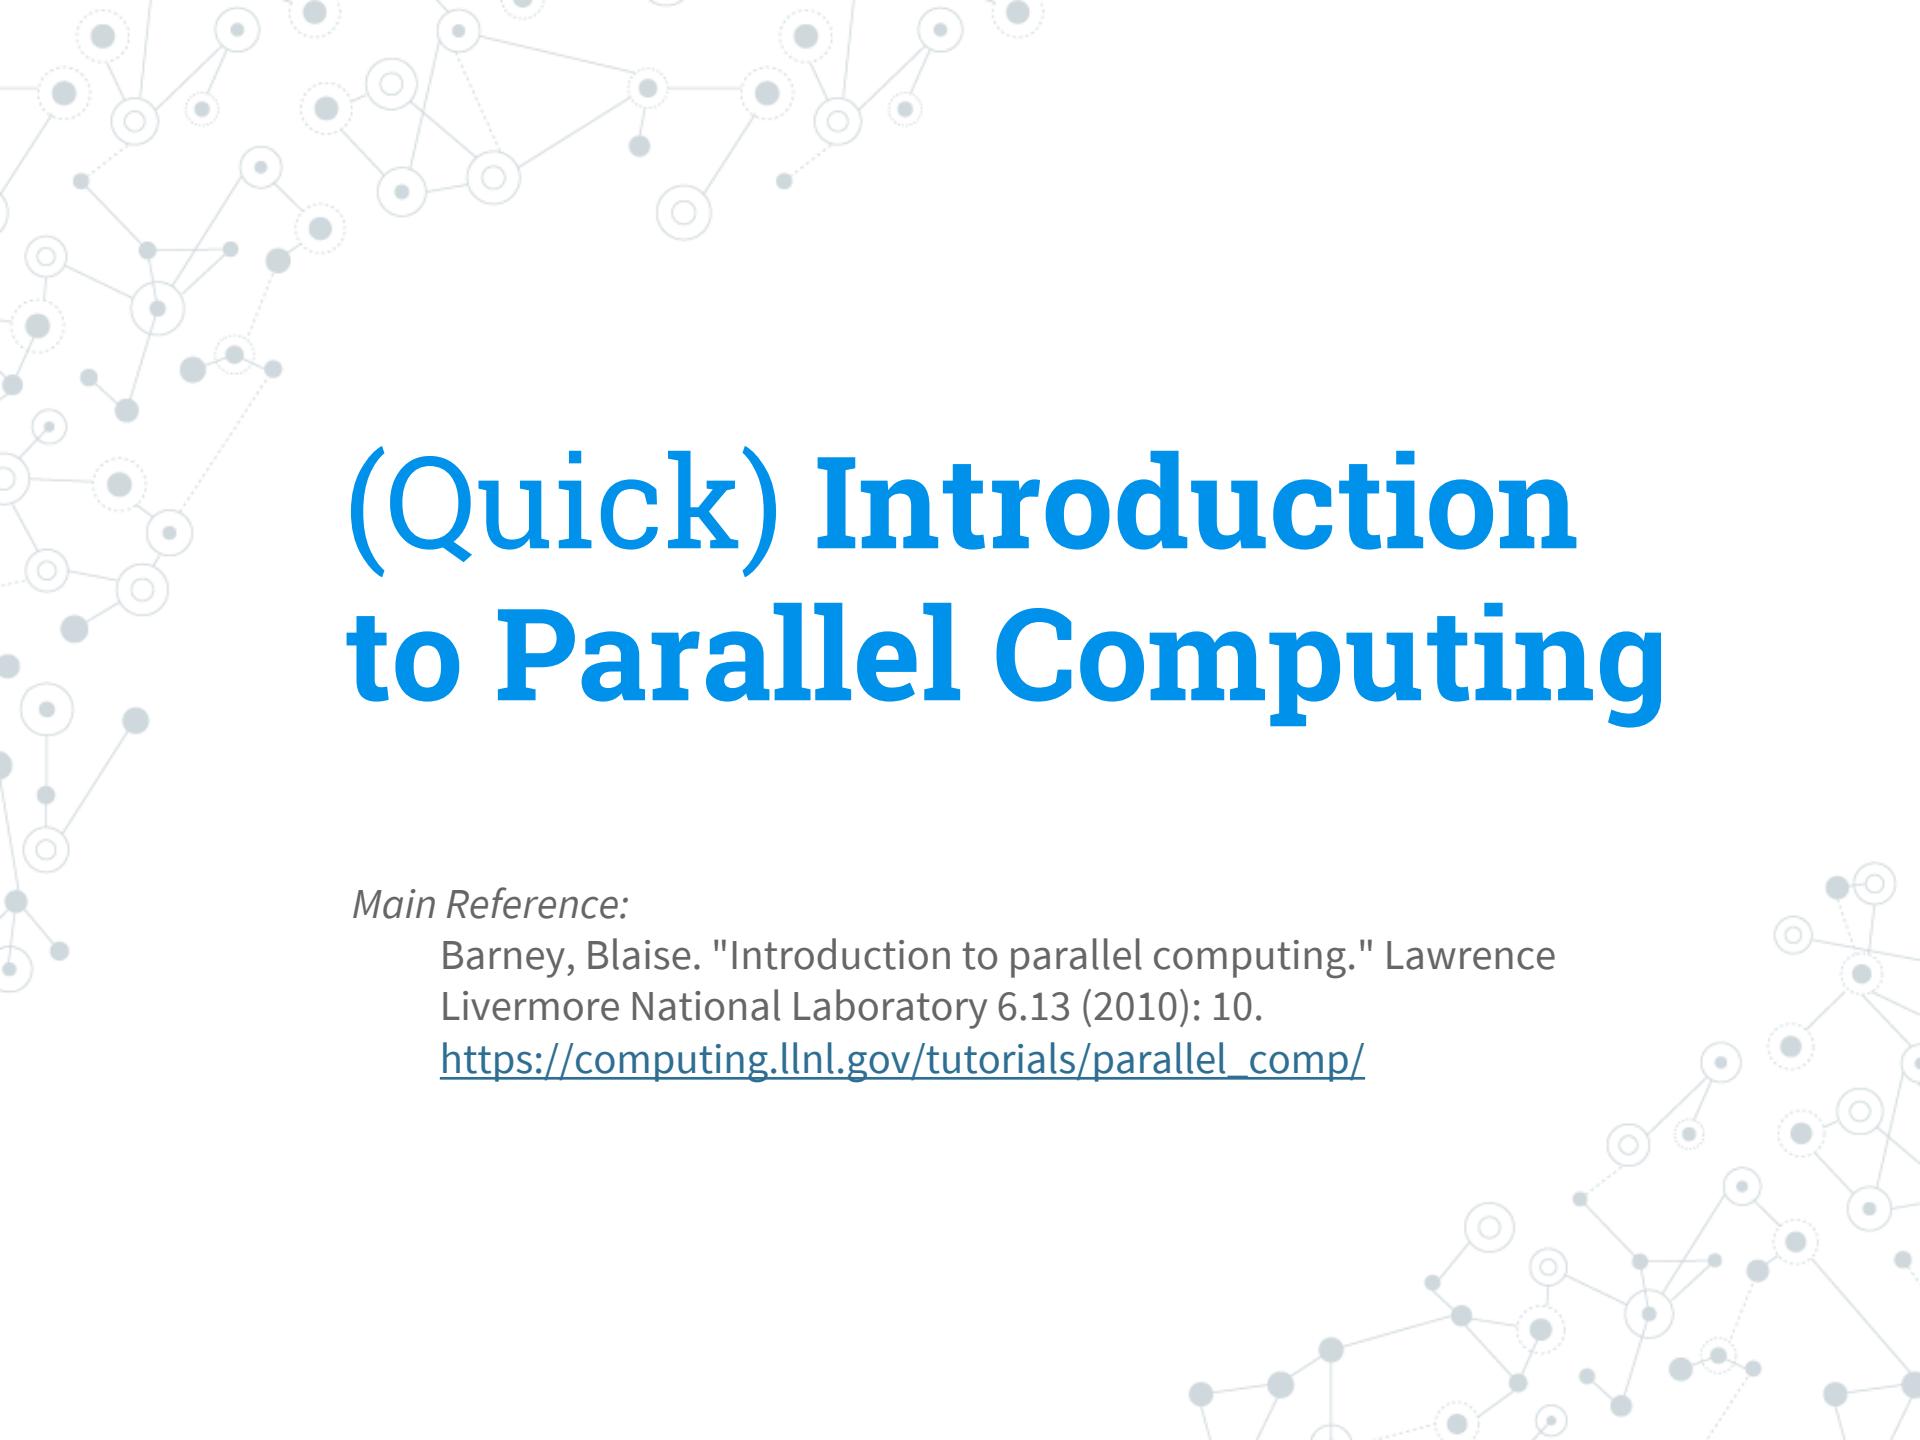
- ④ Developed by J. Garrison and R. Mishmash (UCSB) for the 2013 Trieste summer school on quantum spin liquids
- ④ 1D DMRG code built on numpy and scipy functions:
 - `scipy.sparse.kron` – Kronecker product
 - `scipy.sparse.linalg.eigsh` – Lanczos routine
- ④ Demonstration:
 - `simple_dmrg_02_finite_system.py`
 - `simple_dmrg_03_conserved_quantum_numbers.py`

Features

- ◎ Parallel: Multi-threading capabilities
 - if scipy is compiled with multithreaded BLAS/MKL
- ◎ Built for the 1D Heisenberg XXZ Hamiltonian but it is extensible to other models as well

Limitations

- Ⓐ 2D models often require larger m number of states kept after truncation
- Ⓐ Hamiltonian dimension $\sim m^2 d^2$
 - e.g. for spin-1/2 with $m=2000 \rightarrow 16,000,000$ states
- Ⓐ 2D DMRG requires parallel programming approach



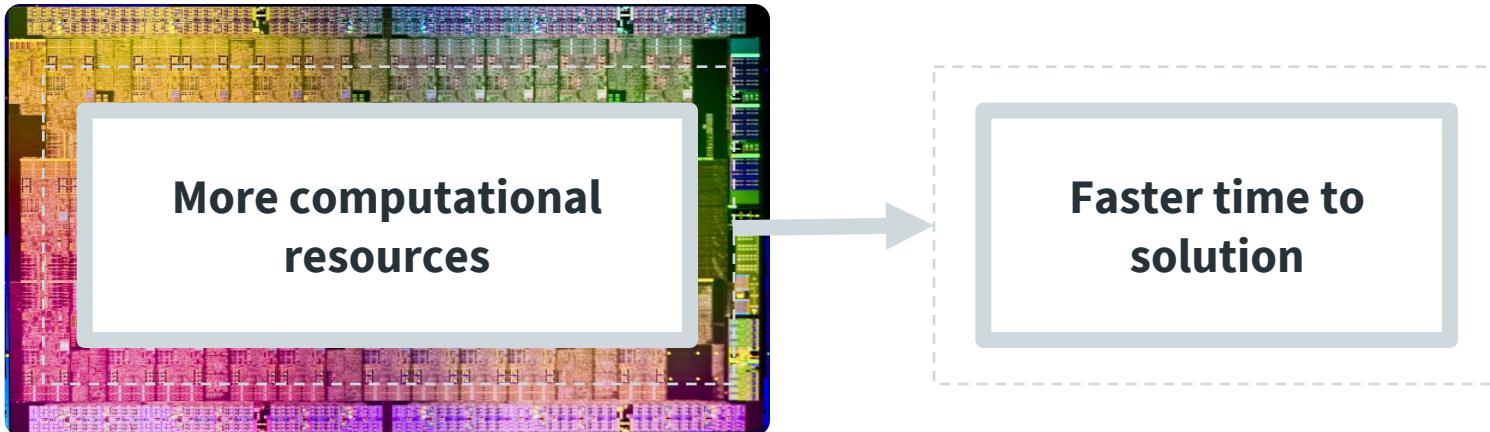
(Quick) Introduction to Parallel Computing

Main Reference:

Barney, Blaise. "Introduction to parallel computing." Lawrence Livermore National Laboratory 6.13 (2010): 10.
https://computing.llnl.gov/tutorials/parallel_comp/

Parallel Computing

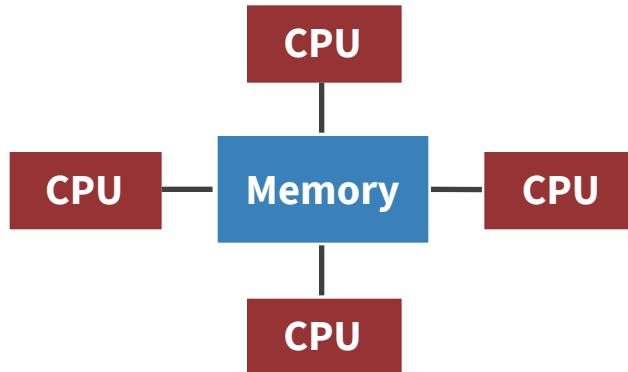
- ◎ A form of computation that performs many calculations simultaneously
- ◎ Instead of making CPUs faster, just **add more cores**
- ◎ Present in most modern computers and mobile devices having “multi-core” technology



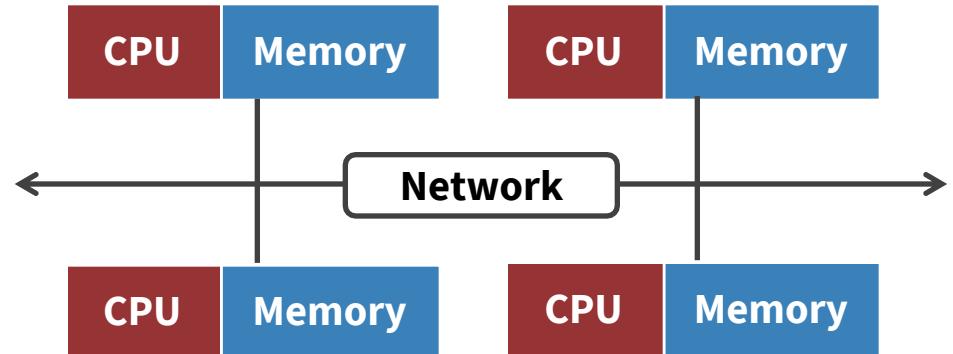
Intel Knights Landing: 72-core Xeon Phi
supercomputing chip

Barney, Blaise. "Introduction to parallel computing." Lawrence Livermore National Laboratory 6.13 (2010): 10. https://computing.llnl.gov/tutorials/parallel_comp/

Parallel Computing Memory Architectures

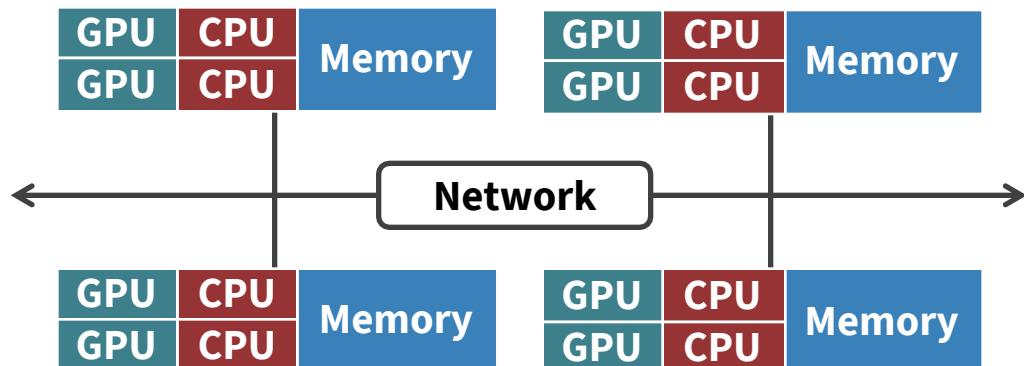


▲ **Shared Memory**
Multithreading, OpenMP



▲ **Distributed Memory**
Message Passing Interface

Hybrid ►
OpenMP-MPI,
CUDA





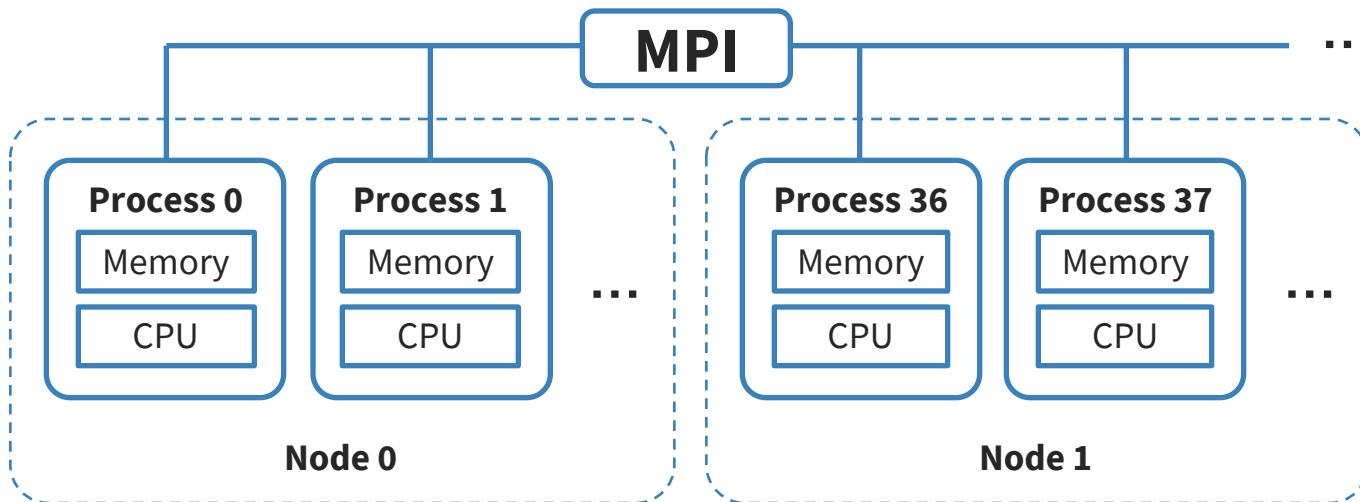
Distributed and Hybrid Systems

Clusters and supercomputers:

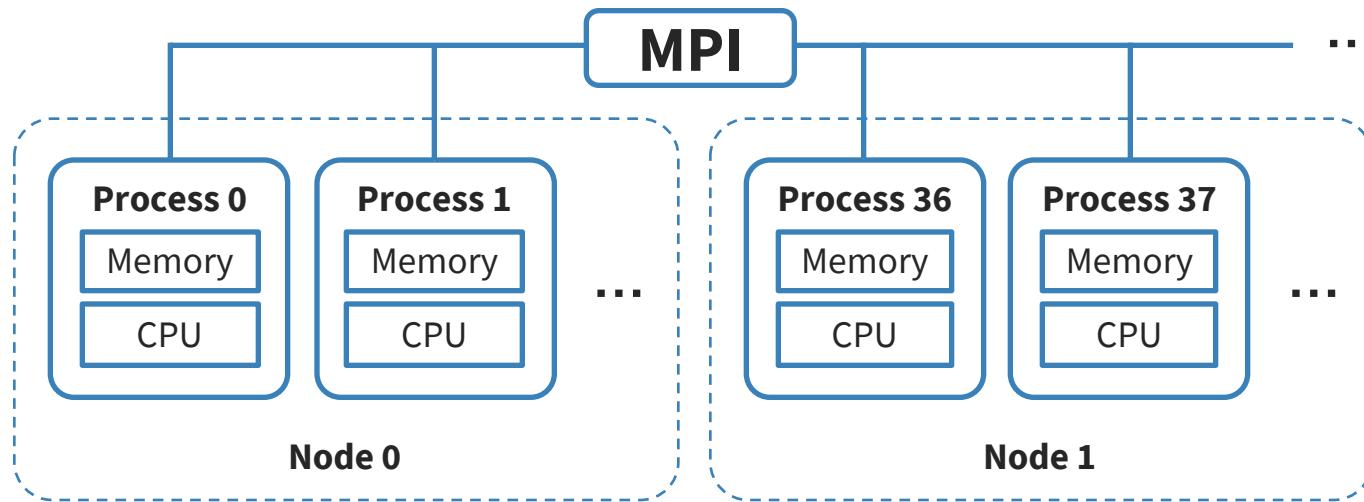
- Argo (ICTP)
- Ulysses (SISSA)
- Marconi (CINECA)

Parallel Programming with Message Passing Interface (MPI)

- MPI is the “de facto” industry standard for distributed systems due to its portability, efficiency, and flexibility
- Runs on *distributed, shared and hybrid* platforms
- Works on a **distributed programming model:**



Parallel Programming with Message Passing Interface (MPI)



- Every process has access only to its own data
- MPI facilitates the communication among processes, whether on the same or different nodes

Some MPI Basics

- ① **MPI communicator** - group of processes that communicate with one another (e.g. `MPI_COMM_WORLD`)
 - ① **rank** – unique id of each process in the communicator
 - ① **size** – total number of processes in the communicator
-
- ① **Types of MPI communication:**
 - Point-to-point
 - Collective

Sample MPI Code

01-hello/hello.c

```
1 #include<mpi.h>
2 #include<stdio.h>
3
4 int main(int argc, char **argv)
5 {
6     int rank, size;
7     MPI_Init(&argc, &argv);
8     MPI_Comm_rank(MPI_COMM_WORLD, &rank);
9     MPI_Comm_size(MPI_COMM_WORLD, &size);
10    /*----- Body of the MPI program -----*/
11
12    printf("Hello from process %d of %d.\n", rank, size);
13
14    /*-----*/
15    MPI_Finalize();
16    return(0);
17 }
```

Compiling the code:

```
$ mpicc hello.c -o hello.x
```

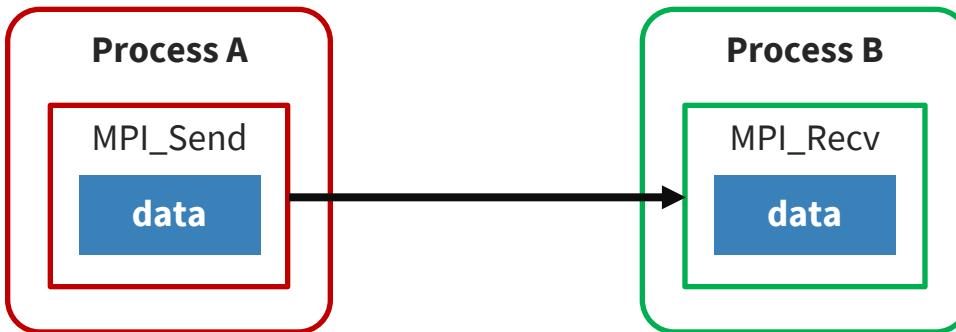
Running the program:

```
$ mpirun -np <NPROCS> <EXEC> <ARGS>
```

e.g.

```
$ mpirun -np 4 ./hello.x
```

Point-to-Point Communication



```
MPI_Send( void* data, int count, MPI_Datatype datatype,  
          int destination, int tag, MPI_Comm communicator)  
MPI_Recv( void* data, int count, MPI_Datatype datatype,  
          int source, int tag, MPI_Comm communicator,  
          MPI_Status* status)
```

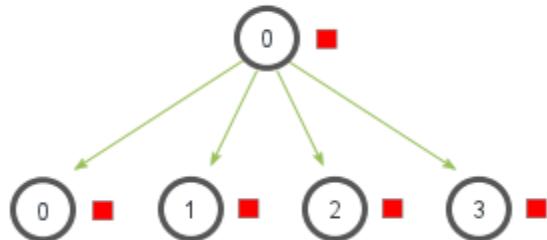
Sample MPI Code

02-send-recv/main.c

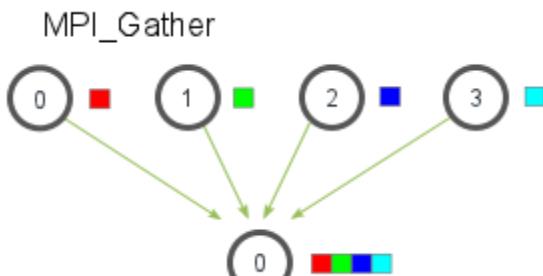
```
10     /*---- Body of the MPI program ----*/
11
12     if(size < 2){printf("Needs at least 2 procs.\n");return(1);}
13
14     int number;
15     if (rank == 0) {
16         number = -1;
17         MPI_Send(&number, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
18     } else if (rank == 1) {
19         MPI_Recv(&number, 1, MPI_INT, 0, 0, MPI_COMM_WORLD,
20                  MPI_STATUS_IGNORE);
21         printf("Process 1 received number %d from process 0\n",
22               number);
23     }
24
25     /*-----*/
```

Collective Communication

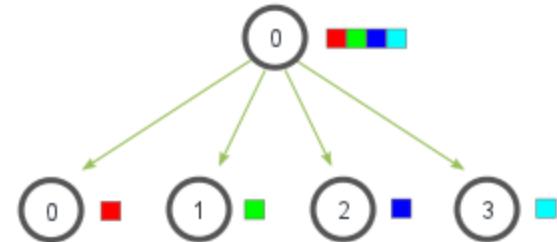
MPI_Bcast



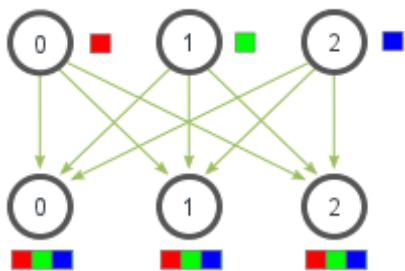
MPI_Gather



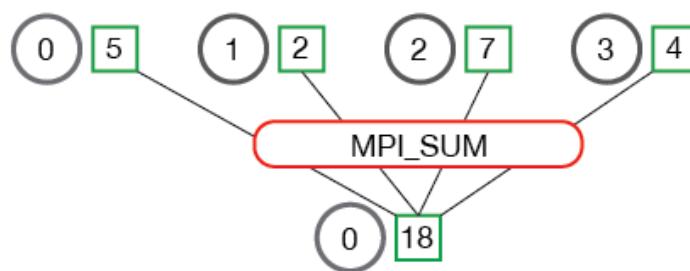
MPI_Scatter



MPI_Allgather



MPI_Reduce



```
MPI_Reduce( void* send_data, void* recv_data, int count,
              MPI_Datatype datatype, MPI_Op op, int root,
              MPI_Comm communicator)
```

Sample MPI Code

03-reduce/main.c

```
10     /*---- Body of the MPI program ----*/
11     /* Each rank r calculates 2*rank+1 */
12     int rank_val = 2*rank + 1;
13     printf("[%d] Val: %-5d\n", rank, rank_val);
14
15     /* Get the sum of all values from all ranks */
16     int sum;
17     MPI_Reduce(&rank_val, &sum, 1, MPI_INT, MPI_SUM, 0,
18                 MPI_COMM_WORLD);
19     printf("[%d] Sum: %-5d\n", rank, sum);
20
21     /* Assert that sum=size^2 */
22     if(rank==0 && sum!=size*size){
23         printf("Sum should be %d.\n", size*size);
24         return(1);
25     }
26     /*-----*/
```

Parallel Linear Algebra

- Recall the Lanczos algorithm

$$|\phi_{n+1}\rangle = \boxed{\hat{H}|\phi_n\rangle} - a_n|\phi_n\rangle - b_n^2|\phi_{n-1}\rangle,$$

$$a_n = \frac{\langle\phi_n|\hat{H}|\phi_n\rangle}{\langle\phi_n|\phi_n\rangle}, \quad b_n^2 = \frac{\langle\phi_n|\phi_n\rangle}{\langle\phi_{n-1}|\phi_{n-1}\rangle},$$

- Large sparse matrix-vector multiplication
- Since the Hamiltonian is very large, we want to parallelize this calculation

Two options:

- Write the parallel algorithms from scratch
- Use third-party libraries

PETSc and SLEPc

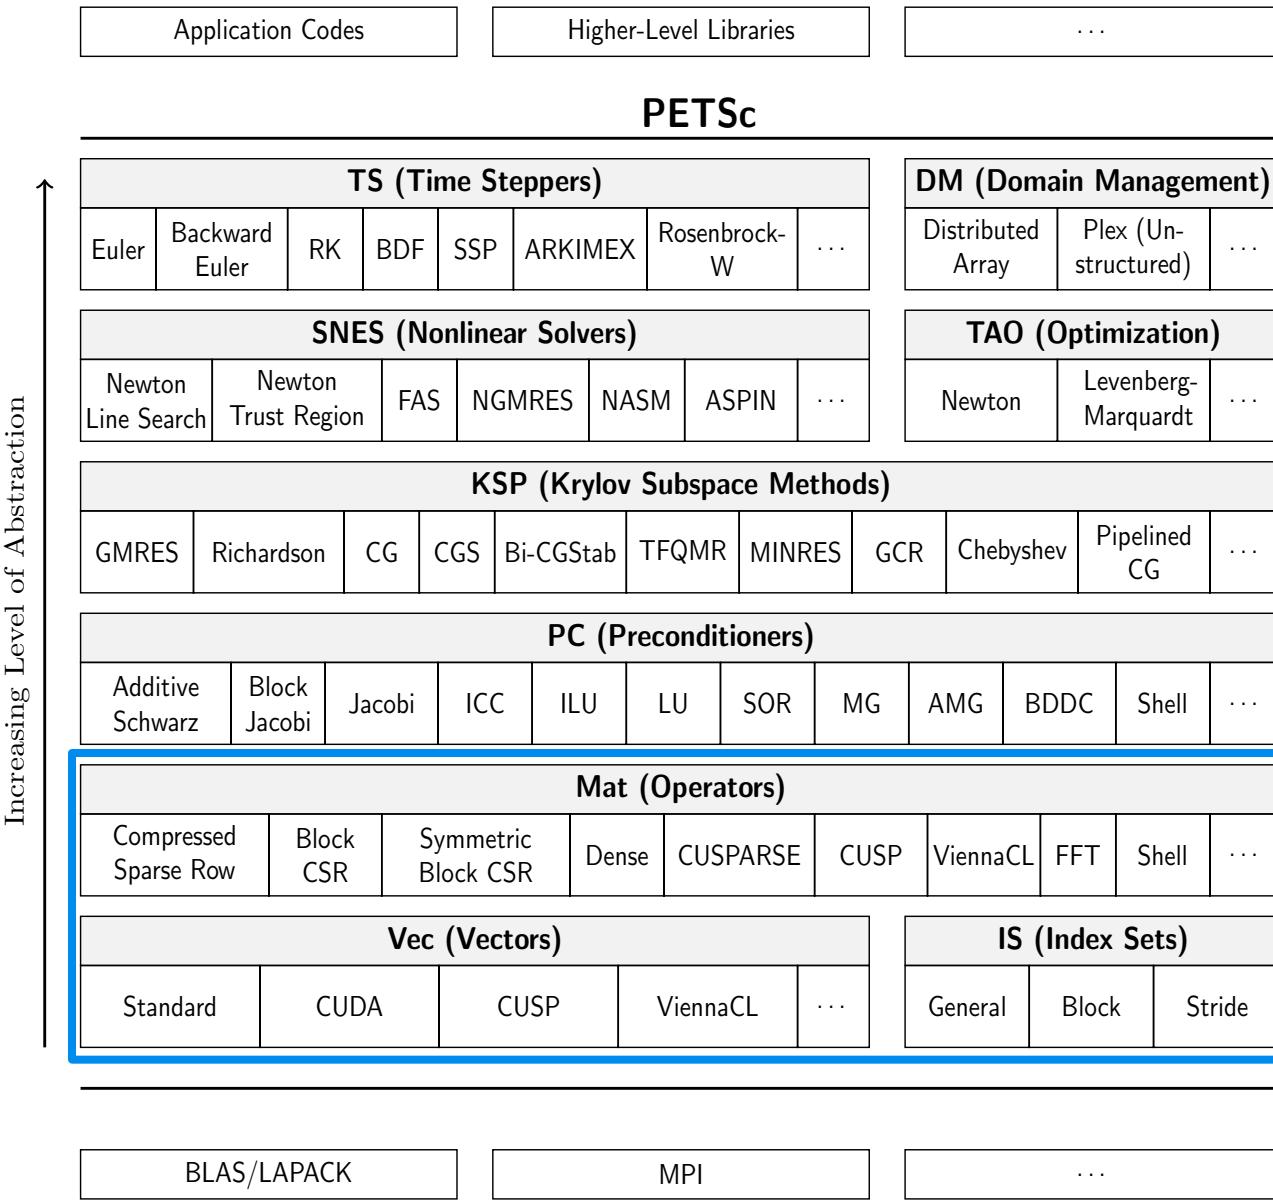
Overview



“

The Portable, Extensible Toolkit for Scientific Computation (PETSc) is a suite of **data structures and routines** for the scalable (parallel) solution of scientific applications modeled by PDEs

Website: www.mcs.anl.gov/petsc/



Numerical libraries of PETSc

A decorative background image featuring a network graph composed of numerous nodes (circles) and edges (lines). A prominent feature is a large, light-gray dashed circle centered on the slide, which contains a blue double quotes symbol (‘’).

The **Scalable Library for Eigenvalue Problem Computations** (SLEPc) is a software library for the solution of large scale **sparse eigenvalue problems** on parallel computers.

Website: slepc.upv.es/

SLEPc

Eigensolvers
for Large
Hamiltonian
Matrices

Nonlinear Eigensolver						M. Function	
SLP	RII	N-Arnoldi	Interp.	CISS	NLEIGS	Krylov	Expokit
Polynomial Eigensolver						SVD Solver	
TOAR	Q-Arnoldi	Linearization		JD	Cross Product	Cyclic Matrix	Thick R. Lanczos
Linear Eigensolver							
Krylov-Schur	Subspace		GD	JD	LOBPCG	CISS	...
Spectral Transformation				BV	DS	RG	FN
Shift	Shift-invert	Cayley	Precond.

Numerical components of SLEPc

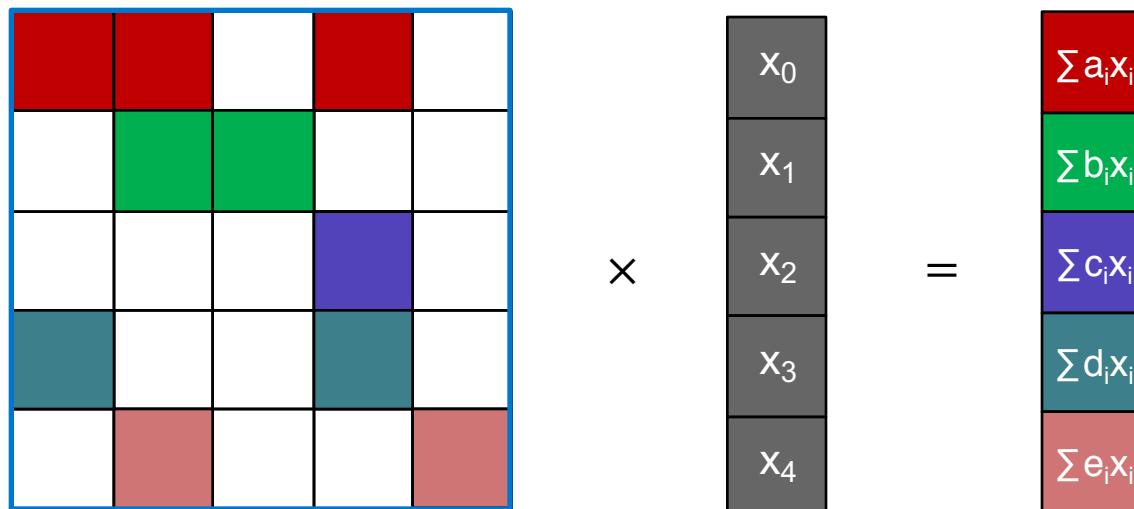
Relevant Features of SLEPc/PETSc

- ◎ Implemented **parallel** data structures and solvers
 - **Vectors** (Vec)
 - **Matrices** (Mat) – Operators and Hamiltonian
 - **Eigenvalue Problem Solver** (EPS) – finding the ground state of the superblock Hamiltonian
- ◎ Portability
 - **Write code once** for different machines, integer and floating point precision, real/complex scalars
 - Interfaces to BLAS/LAPACK routines
- ◎ Parallel Efficiency
 - Implemented algorithms are already tried and tested for efficiency (no need to code from scratch)

Parallel Linear Algebra with PETSc

◎ Sparse Matrix-Vector Multiplication

Logical layout:

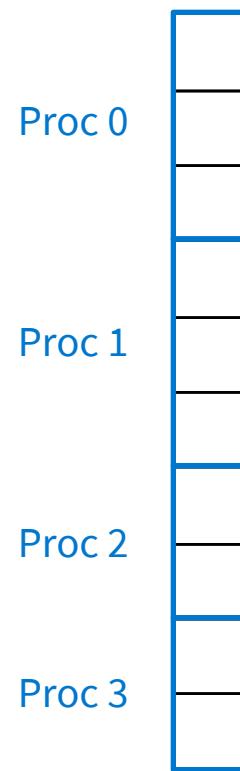


Compressed sparse row (CSR):

(0,a ₀)	(1,a ₁)	(3,a ₃)	(1,b ₁)	(2,b ₂)	(3,c ₃)	(0,d ₀)	(3,d ₃)	(1,e ₁)	(4,e ₄)
---------------------	---------------------	---------------------	---------------------	---------------------	---------------------	---------------------	---------------------	---------------------	---------------------

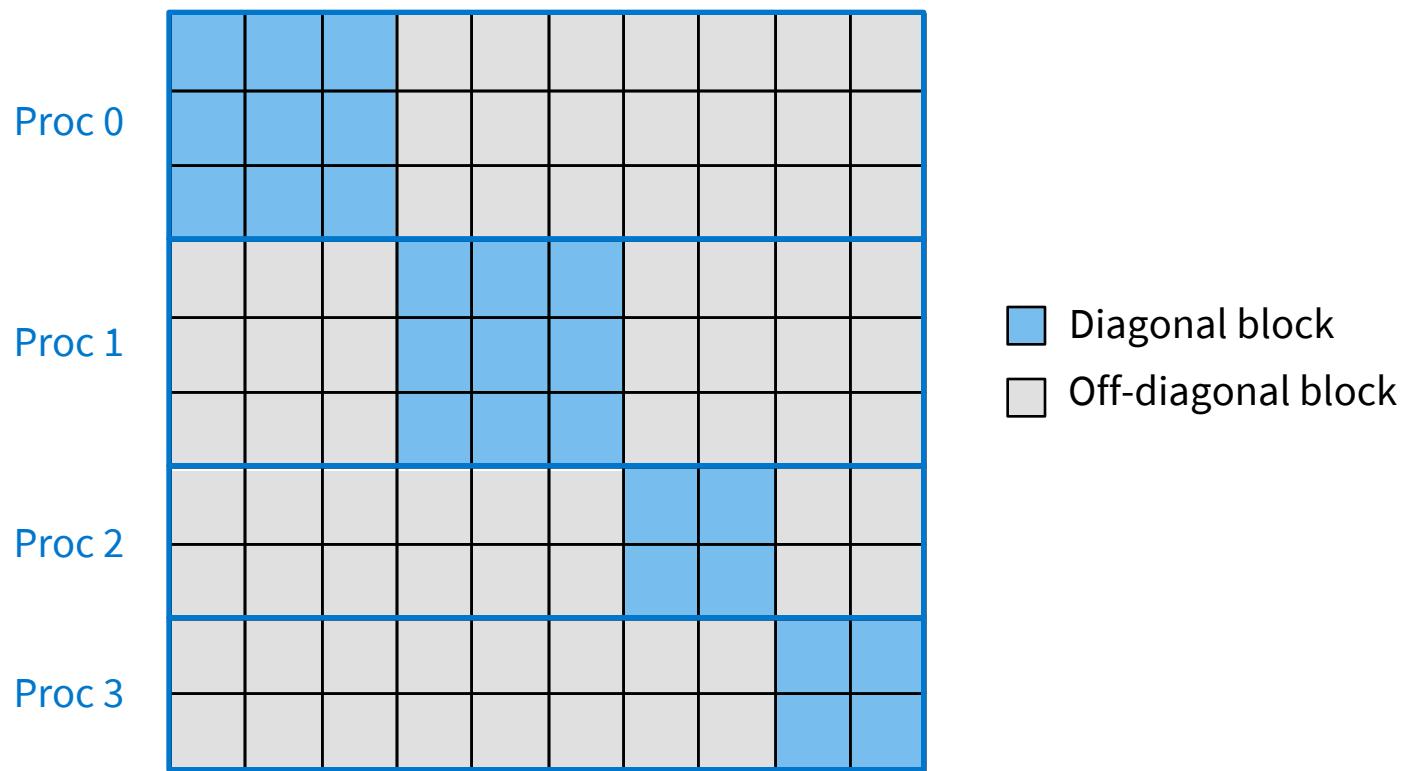
Parallel Linear Algebra with PETSc

⌚ Distributed Vectors



Parallel Linear Algebra with PETSc

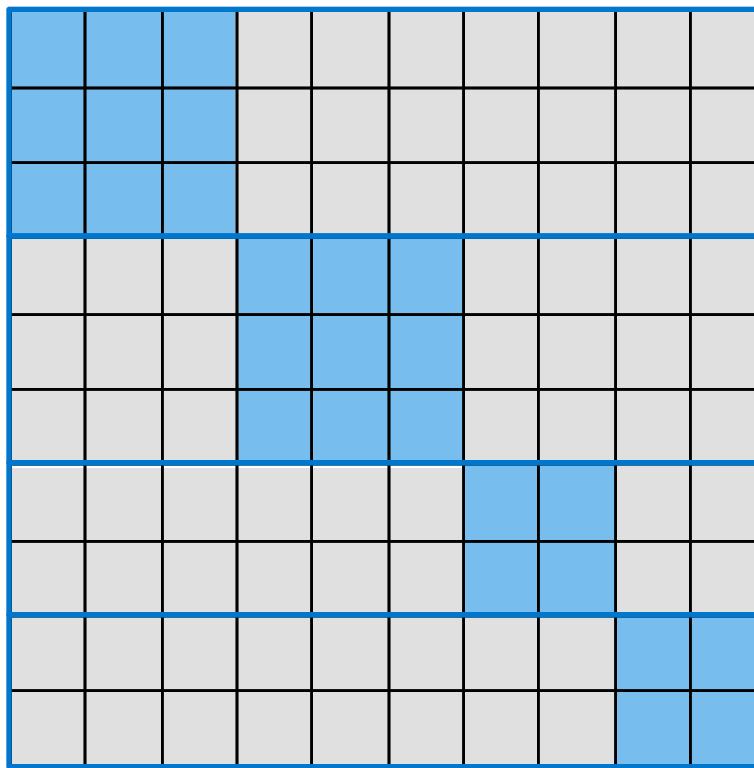
◎ Row-Distributed Matrices



Parallel Linear Algebra with PETSc



Parallel Sparse Matrix-Vector Multiplication

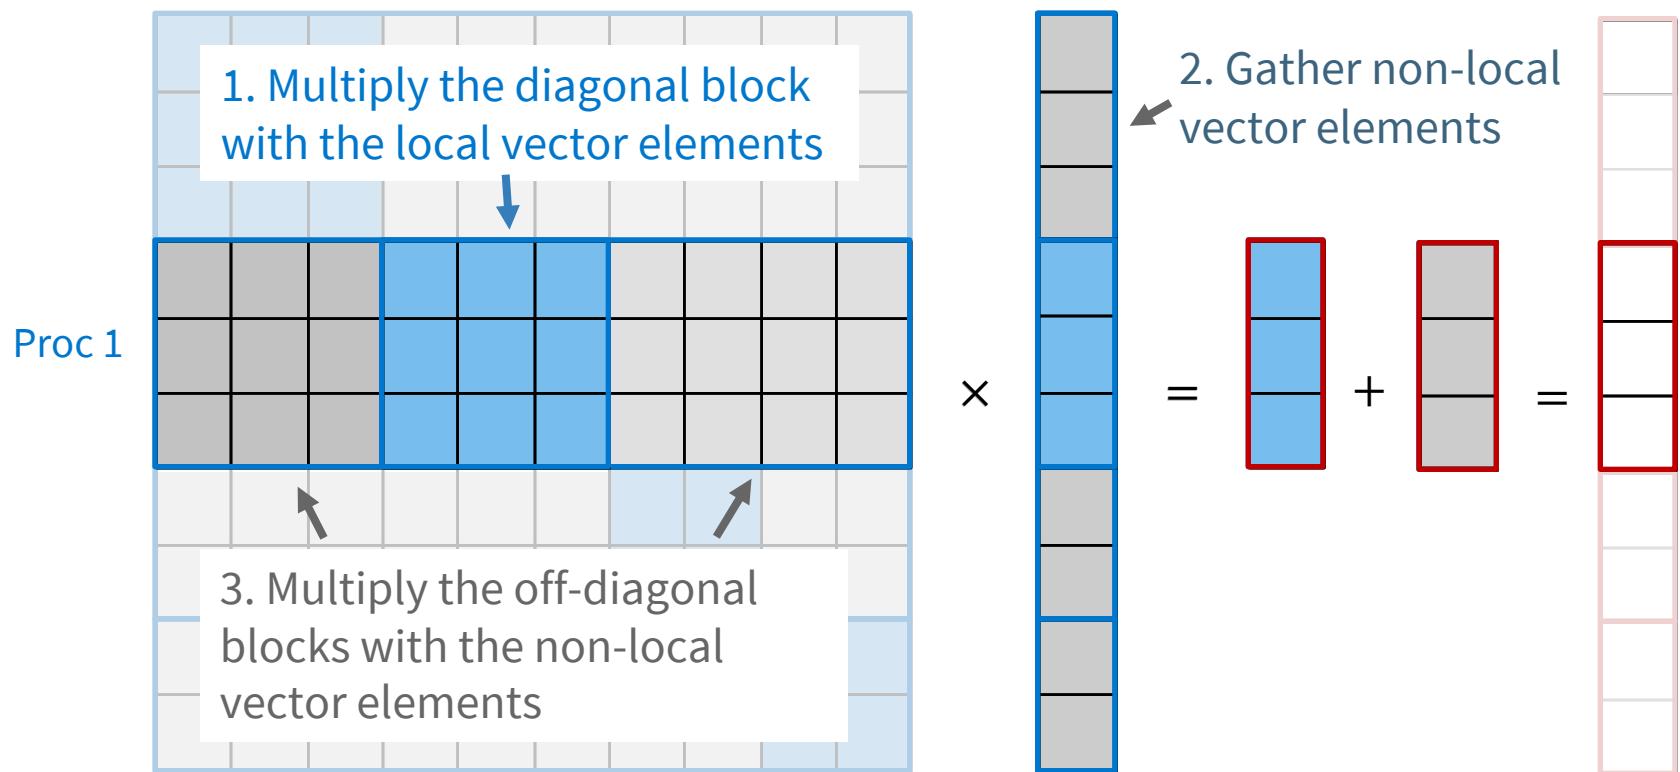


CSR matrix blocks

Parallel Linear Algebra with PETSc



Parallel Sparse Matrix-Vector Multiplication



Finding the Ground State with SLEPc

◎ Iterative solution using Krylov subspace methods

$$\mathcal{K}_m(A, v) \equiv \text{span} \{v, Av, A^2v, \dots, A^{m-1}v\}$$

◎ Built-in solvers:

- Arnoldi
- Lanczos
- **Krylov-Schur**
- etc.

For Hermitian matrices, it is equivalent to the **thick-restarted Lanczos algorithm** (keeping vectors orthogonal to machine ϵ)

Installing the libraries

- Ⓐ A step-by-step tutorial can be found at:

<http://dmrgx.rtfd.io>

under **Related Pages** > **Installation**

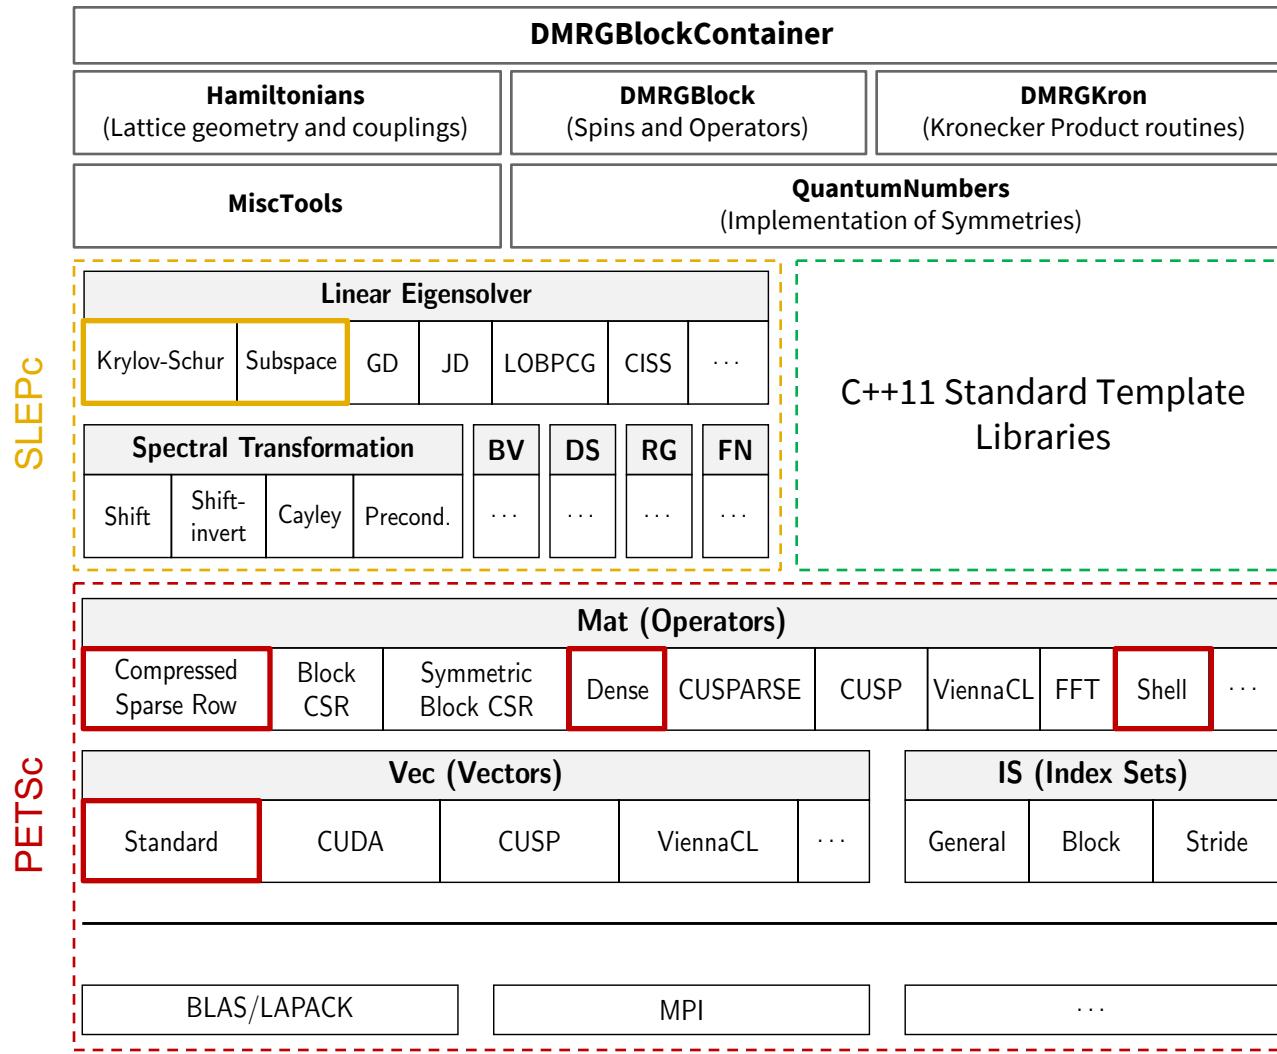
PETSc/SLEPc Exercises

- Ⓐ Hello World
- Ⓑ SLEPc Hands-on 1

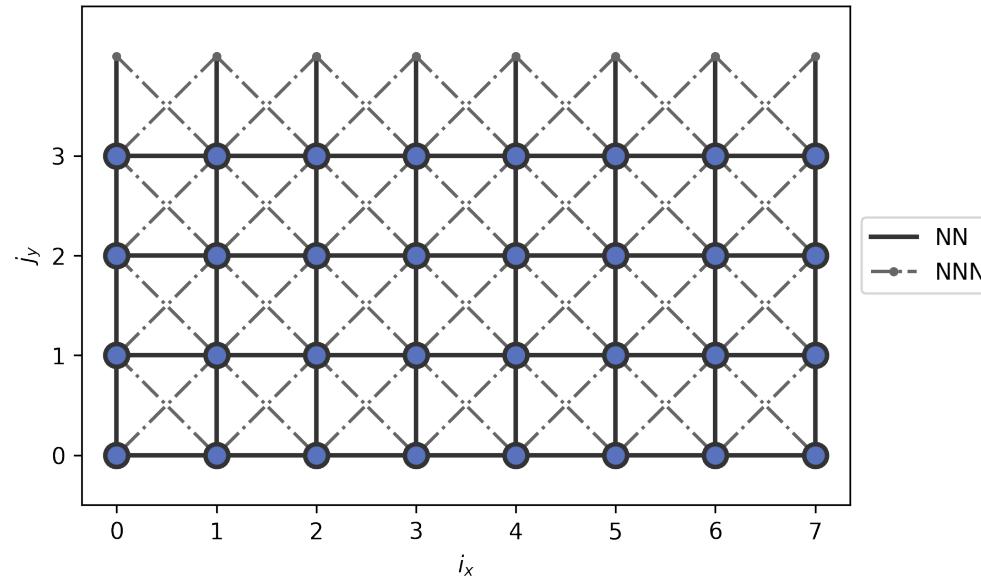
see SampleCode/02-petsc-slepc

DMRG.x

DMRG.x Modules and Libraries



Hamiltonian and Geometry



$$H = \sum_{\langle ij \rangle} \left(J_1 \left[\hat{S}_i^+ \hat{S}_j^- + \hat{S}_i^- \hat{S}_j^+ \right] + \Delta_1 \hat{S}_i^z \hat{S}_j^z \right)$$
$$+ \sum_{\langle\langle kl \rangle\rangle} \left(J_2 \left[\hat{S}_k^+ \hat{S}_l^- + \hat{S}_k^- \hat{S}_l^+ \right] + \Delta_2 \hat{S}_k^z \hat{S}_l^z \right)$$

2D-DMRG

Using a one-dimensional traversal

see 2D-DMRG.ppt

Implementation of U(1) Symmetries

- ◎ Since the Hamiltonian conserves total magnetization, we can look for the ground state at a specific total magnetization sector.
- ◎ The sectors are stored as lookup keys that map to a range of indices containing that array

$$S_j^z \rightarrow [i_{\text{start}}, i_{\text{end}} + 1)$$

- ◎ Then the Hamiltonian is built using sectors that satisfy

$$S_{j,A}^z + S_{k,B}^z = S_{\text{tot}}^z$$

Usage

- Ⓐ A step-by-step tutorial can be found at:

<http://dmrgx.rtfd.io>

under **Related Pages > Usage**

Post-Processing

Post-processing module written in Python

◎ DMRG.x/postproc/dmrg_postprocessing.py

Recommended interface: **jupyter notebook**

Credits

Special thanks to all the people who made and released these awesome resources for free:

- Presentation template by [SlidesCarnival](#)
- Photographs by [Unsplash](#) & [Death to the Stock Photo license](#)

**More info on how to use this template at
www.slidescarnival.com/help-use-presentation-template**

This template is free to use under [Creative Commons Attribution license](#). You can keep the Credits slide or mention SlidesCarnival and other resources used in a slide footer.